

An Introduction to Apache Spark

Performance Tuning

Different parts of Spark Job to optimize

- Code-level design choices (e.g., RDDs versus DataFrames)
- Data at rest
- Joins
- Aggregations
- Data in flight
- Individual application properties
- Inside of the Java Virtual Machine (JVM) of an executor
- Worker nodes
- Cluster and deployment properties

Implement monitoring and job history tracking

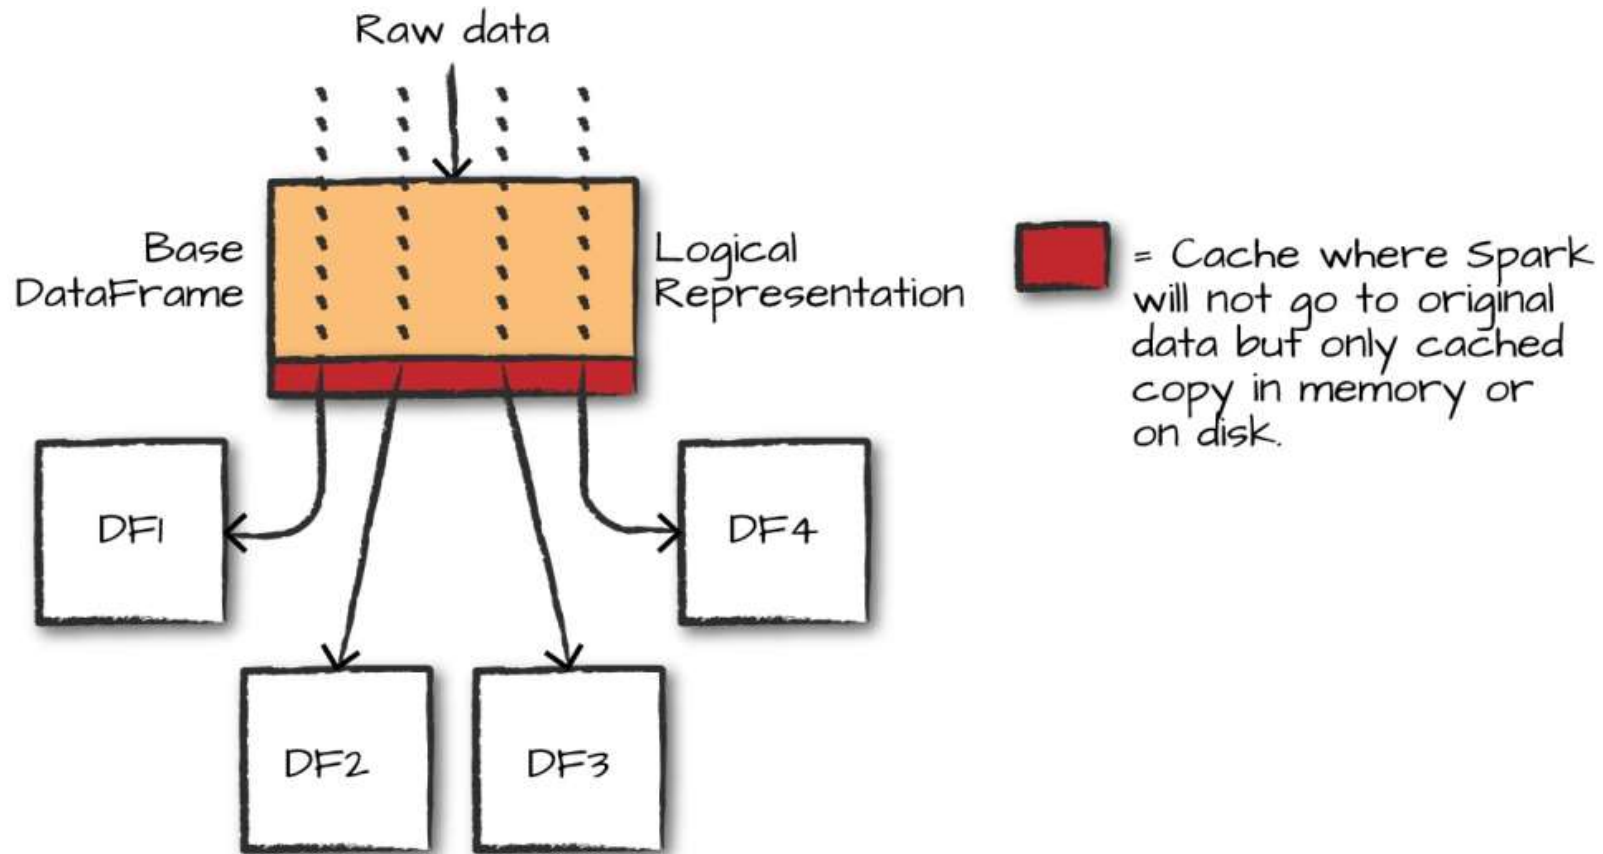
- To figure out how to improve performance
- To know whether you're really improving job performance.

Indirect Performance Enhancements

- Design Choices
 - Scala versus Java versus Python versus R
 - Depending on the use case
- Data at Rest
 - Making sure that you're storing your data for effective reads later on is absolutely essential to successful big data projects.
- Caching
 - Will place a DataFrame, table, or RDD into temporary storage across the executors in your cluster, and make subsequent reads faster

Caching

- We can avoid having to recompute the original DataFrame (i.e., load and parse the CSV file) many times by adding a line to cache



Caching

- `DF1 = spark.read.format("csv")\`
- `.option("inferSchema", "true")\`
- `.option("header", "true")\`
- `.load("/data/flight-data/csv/2015-summary.csv")`
- `DF2 = DF1.groupBy("DEST_COUNTRY_NAME").count().collect()`
- `DF3 = DF1.groupBy("ORIGIN_COUNTRY_NAME").count().collect()`
- `DF4 = DF1.groupBy("count").count().collect()`

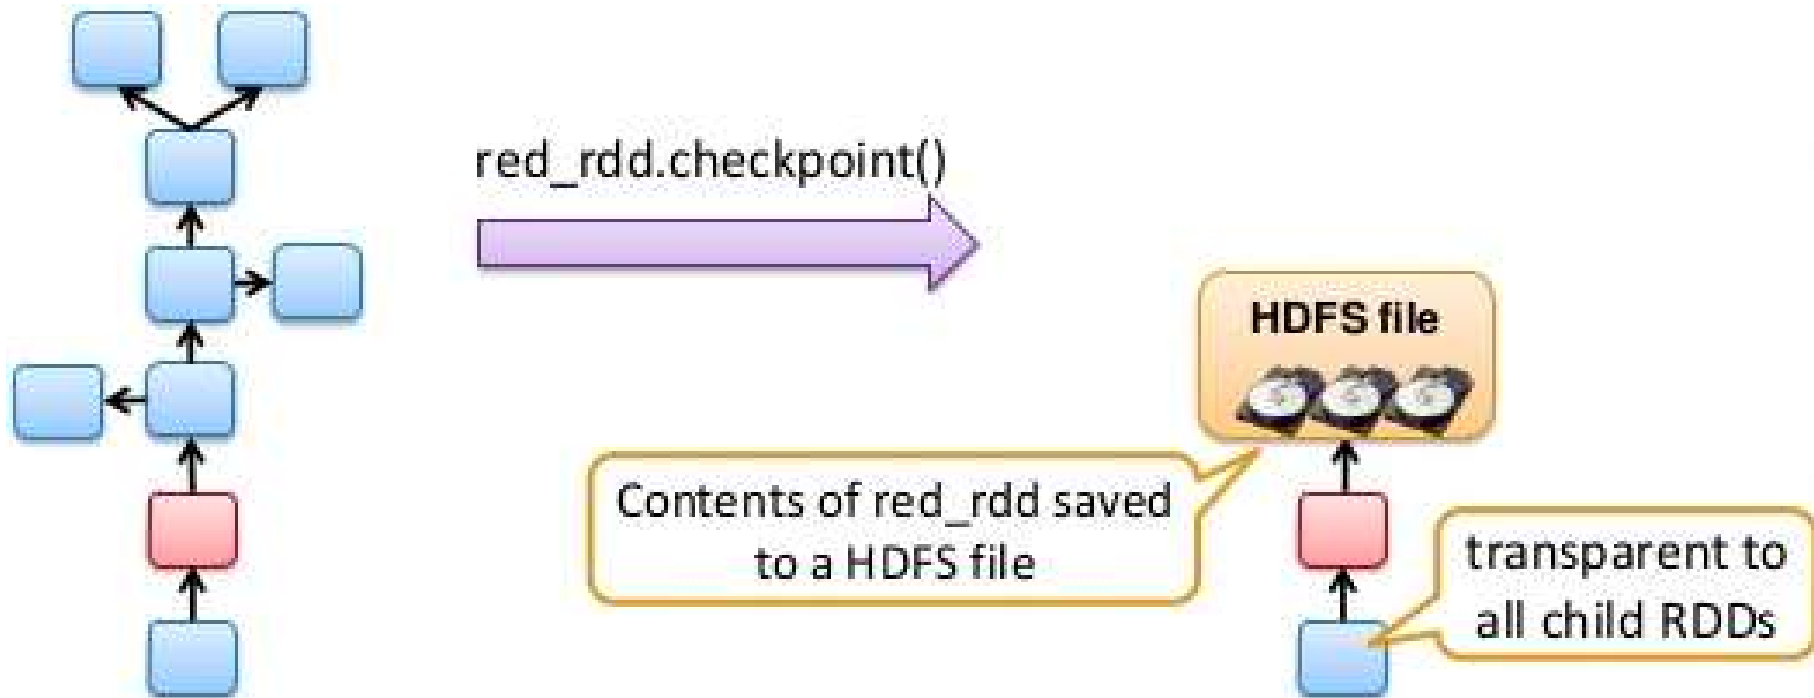
- `DF1.cache()`
- `DF1.count()`

Caching

- Now that the data is cached, the commands will be faster, as we can see by running the following code:
 - `DF2 = DF1.groupBy("DEST_COUNTRY_NAME").count().collect()`
 - `DF3 = DF1.groupBy("ORIGIN_COUNTRY_NAME").count().collect()`
 - `DF4 = DF1.groupBy("count").count().collect()`

What is Check-pointing?

- Saving RDD to HDFS to prevent RDD graph from growing too large
- Will persist the transformed RDD or DataFrame forever.




When should I cache or checkpoint?

- Determine if the results of a set of transformations can be reused for a very long time or not
 - If the answer is yes, use **checkpointing**
 - If the answer is no, use **caching**
- Example when checkpointing would be preferred:
 - Crunching a RDD or DataFrame of taxes for a previous year
 - They are unlikely to change once calculated so it would be much better to checkpoint and save them forever

Shuffling


Narrow transformation

- Input and output stays in same partition
- No data movement is needed



Wide transformation

- Input from other partitions are required
- Data shuffling is needed before processing



Minimizing Shuffling for Increased Performance

- Shuffle is an expensive operation
- Each shuffling generates a new stage.
- Here are some tips to reduce shuffle:
 - Use the Spark UI to study the plan to look for opportunity to reduce the shuffle
 - Use the built in `aggregateByKey()` operator instead of writing your own aggregations.
 - Filter input earlier in the program rather than later.
 - repartition, join, cogroup, and any of the `*By` or `*ByKey` transformations can result in shuffles.

Thanks